

```

import numpy as np
from tqdm import tqdm

class Linear_Regression():
    def __init__(self, alpha = 1e-3 , num_iter = 10000, early_stop = 100,
        intercept = True, init_weight = None, penalty = None,
        lam = 1e-3, normalize = False, adaptive=True):
        """
        Linear Regression with gradient descent method.

        Attributes:
        -----
        alpha: Learning rate.
        num_iter: Number of iterations
        early_stop: Number of steps without improvements
                    that triggers a stop training signal
        intercept: True = with intercept (bias), False otherwise
        init_weight: Optional. The initial weights passed into the model,
                    for debugging
        penalty: {None,l1,l2}. Define regularization type for regression.
                None: Linear Regression
                l1: Lasso Regression
                l2: Ridge Regression
        lam: regularization constant.
        normalize: True = normalize data, False otherwise
        adaptive: True = adaptive learning rate, False = fixed learning rate
        """
        self.alpha = alpha
        self.num_iter = num_iter
        self.early_stop = early_stop
        self.intercept = intercept
        self.init_weight = init_weight
        self.penalty = penalty
        self.lam = lam
        self.normalize = normalize
        self.adaptive = adaptive

    def fit(self, X, y):
        # initialize X, y
        self.X = X
        self.y = np.array([y]).T

        self.max = np.zeros((1, X.shape[1]))
        self.min = np.zeros((1, X.shape[1]))
        for i in range(self.X.shape[1]):
            self.max[:, i] = self.X[:, i].max()
            self.min[:, i] = self.X[:, i].min()

        ##### START TODO 1 #####
        # Normalize the data using the formula provided in lecture
        if self.normalize:
            self.X = (self.X - self.min) / (self.max - self.min)
        ##### END TODO 1 #####

        ##### START TODO 2 #####
        # Add bias (if necessary) by concatenating a constant column into X
        # Hint: go through HW1 Q5 might be helpful
        if self.intercept:
            col = np.ones((len(self.X), 1))
            self.X = np.hstack((self.X, col))
        ##### END TODO 2 #####

        # initialize coefficient
        self.coef = self.init_weight if self.init_weight is not None\
        else np.array([np.random.uniform(-1,1,self.X.shape[1])]).T

        # start training, self.loss is used to record losses over iterations
        self.loss = []
        self.gradient_descent()

    def gradient(self):
        coef = -2 / len(self.X)

        ##### START TODO 3 #####
        # Find prediction and gradient

```

```

# Hint: Find the model's prediction from the given inputs with the
#       coefficient, then calculate the gradient
# If you forgot the formula, find them in lecture 4 and 5
pred = self.X @ self.coef
grad = coef * (self.X.T @ (self.y - pred))
##### END TODO 3 #####

##### START TODO 4 #####
# Implement regularization penalty
# Hint: Use self.lam
if self.penalty == 'l2':
    grad += 2 * self.lam * self.coef
elif self.penalty == 'l1':
    grad += self.lam * np.sign(self.coef)
else:
    pass
##### END TODO 4 #####

return grad

def gradient_descent(self):
    print('Start Training')
    for i in range(self.num_iter):

        ##### START TODO 5 #####
        # calculate prediction y based on current coefficients (self.coef)
        previous_y_hat = self.X @ self.coef
        grad = self.gradient()
        # calculate the new coefficients after incorporating the gradient
        temp_coef = self.coef - (self.alpha * grad)
        ##### END TODO 5 #####

        ##### START TODO 6 #####
        # calculate regularization cost (alias: regularization loss) based on
        # self.coef and temp_coef
        if self.penalty == 'l2':
            previous_reg_cost = self.lam * np.sum(np.square(self.coef))
            current_reg_cost = self.lam * np.sum(np.square(temp_coef))
        elif self.penalty == 'l1':
            previous_reg_cost = self.lam * np.sum(np.abs(self.coef))
            current_reg_cost = self.lam * np.sum(np.abs(temp_coef))
        else:
            previous_reg_cost = 0
            current_reg_cost = 0
        ##### END TODO 6 #####

        ##### START TODO 7 #####
        # Calculate error (alias: loss) using sum squared loss
        # and add regularization cost
        pre_error = np.sum(np.square(self.y - previous_y_hat)) + previous_reg_cost
        current_y_hat = self.X @ temp_coef
        current_error = np.sum(np.square(self.y - current_y_hat)) + current_reg_cost
        ##### END TODO 7 #####

        # Early Stop: early stop is triggered if loss is not decreasing
        # for some number of iterations
        if len(self.loss) > self.early_stop and \
        self.loss[-1] >= max(self.loss[-self.early_stop:]):
            print('-----')
            print(f'End Training (Early Stopped at iteration {i})')
            return self

        ##### START TODO 8 #####
        # Implement adaptive learning rate

        # Rules: if current error is smaller than previous error,
        # multiply the current learning rate by 1.3 and update coefficients,
        # otherwise by 0.9 and do nothing with coefficients
        if current_error < pre_error:
            self.alpha = 1.3 * self.alpha if self.adaptive else self.alpha
            self.coef = temp_coef
        else:
            self.alpha = 0.9 * self.alpha if self.adaptive else self.alpha
        ##### END TODO 8 #####

```

```

        # record stats
        self.loss.append(float(current_error))
        if i % 1000000 == 0:
            print('-----')
            print('Iteration: ' + str(i))
            print('Coef: ' + str(self.coef))
            print('Loss: ' + str(current_error))
        print('-----')
    print('End Training')
    return self

def predict(self, X):
    X_norm = np.zeros(X.shape)
    for i in range(X.shape[1]):
        X_norm[:, i] = (X[:, i] - self.min[:, i]) / (self.max[:, i] - self.min[:, i])
    X = X_norm
    ##### START TODO 9 #####
    # add bias (if necessary, same as TODO 2)
    if self.intercept:
        col = np.ones((len(X), 1))
        X = np.hstack((X, col))

    # Find the model's predictions
    # Hint: Use matrix multiplication ('@' might come in handy here)
    y = X @ self.coef
    return y
    ##### END TODO 9 #####

# Congrats! You have reached the end of this model's implementation :)

```